

Active sensing is not communication

Causal auditing of emergent signals in embodied multi-agent reinforcement learning, and the channel factorisation that lets signalling appear

Abstract

Every electric organ discharge (EOD) a weakly electric fish produces does two things at once: it illuminates the emitter’s own surroundings, and it broadcasts the emitter’s presence to anyone with the receptors to detect it. This dual use makes it unusually hard to say whether an EOD is *communication*. A recent multi-agent reinforcement learning model of electric fish collectives showed that EODs shape group foraging and spacing, and defined communication operationally as an EOD that conspecifics detect and respond to. That definition cannot separate a signal from a cue, and the single intervention available in that model—silencing one fish—removes the emitter’s own electrolocation and its neighbours’ access to it simultaneously, so the resulting behavioural change cannot be attributed to either. We build a GPU-vectorised reimplementation of that simulator, validated against the original to floating-point round-off (maximum relative deviation 6×10^{-10} across all sensory pathways), that runs 596k agent-steps per second and lets us switch the consequences of a discharge independently. A single motor act is decomposed into three separately controllable channels: *reafference* (the emitter’s own image of the world), *illumination* (the same pulse lighting up a neighbour’s world—an exploitable cue), and *detection* (the pulse registering on a neighbour’s knollenorgans—a candidate signal). We further add what the original model lacks: a metabolic cost per discharge, electroreceptive predators that home in on EODs, switchable cooperative and competitive food economics, tunable social range, and persistent individual identity. We first validate the assays on dyads that are communicating by construction: in a referential game embedded in the same physics they separate informative from uninformative senders and attentive from deaf receivers, and deleting a genuinely used signal costs its receiver -172.5 in return. Each assay alone, however, is foolable — positive signalling is just as high when nobody listens, and causal influence just as high when the channel carries nothing — so only their conjunction identifies communication. Applied to the electric discharge, silencing a fish costs it -0.52 food items per episode, of which -0.51 is the loss of its own electrolocation and 0.03 the loss of its detectability by others; the private share stays between 83% and 135% across group sizes, arena scales, episode lengths and network widths. The public channel satisfies both standard criteria — decodable information about food ($\Delta R^2 = 0.156$) and receiver-policy influence $13\times$ the noise floor — yet replaying, scrambling or fabricating it moves no payoff we can resolve, sender shaping is 0.042 once reception is held fixed and only audibility is cut, and metabolic cost and eavesdropping predation regulate discharge rate while making the pulse train *less* informative. Freeing one bit of discharge subtype from the sensory function does not change this. A system can pass every standard diagnostic for emergent communication and not be communicating; what the discharge lacks is not a channel but a task in which its private information is worth anything to anyone else. Standard emergent-communication diagnostics can therefore misclassify an exploited sensory cue as communication, and freeing the channel is not by itself what converts one into the other.

1 Introduction

A weakly electric fish navigates by emitting brief electric pulses and reading the distortions that nearby objects impose on the resulting field [1, 2]. The same pulses propagate far beyond the range at which they are useful for electrolocation, and other fish—equipped with dedicated knollenorgan electroreceptors tuned to conspecific discharges—can detect them at up to a metre [3]. An EOD is therefore simultaneously a private sensory probe and a public event. Which of these functions explains why the fish emits when it does is the central question of electric-fish neuroethology, and it is an instance of a much older problem in animal communication: distinguishing a *signal*, whose production has been shaped by its effect on receivers, from a *cue*, which receivers merely exploit [4, 5].

Singh et al. [6] introduced a computational framework in which agents with biophysically grounded electrosensory and electromotor systems are trained by multi-agent reinforcement learning to forage collectively. Their agents reproduce several features of real fish, and interventions on the model—ablating electroreceptor classes, silencing an agent’s EOD, and altering the food distribution—show that discharges influence group foraging and social spacing. They define communication operationally, as “EOD signals that are sensed by conspecifics via Knollenorgans and that influence their behavior.”

That definition is a statement about effects, and effects are cheap. The emergent-communication literature has repeatedly shown that behavioural association between one agent’s output and another’s behaviour is not sufficient evidence of signalling: a receiver may be responding to something correlated with the message, and a sender’s output may be informative without having been shaped by that informativeness at all [7, 8]. Lowe et al. [7] formalise the two requirements as *positive signalling* (the message carries information about the sender’s state) and *positive listening* (the message causally changes the receiver’s behaviour), and note that either can occur without the other.

For electric fish the problem is sharper than in a generic emergent-communication setup, because the channel is not a dedicated cheap-talk symbol but a physical act with an unavoidable private benefit. Three consequences follow from a single discharge:

1. **Reafference.** The pulse casts an electric image of nearby objects onto the emitter’s own mormyromasts. This is active electrolocation and is the reason the behaviour exists at all.
2. **Illumination.** The same field polarises objects in a neighbour’s vicinity, and the neighbour’s mormyromasts can read those images—collective electrosensing, demonstrated in real fish by Pedraja and Sawtell [9]. The neighbour gains information from a probe it did not pay for. This is the paradigm case of a *cue*.
3. **Detection.** The pulse registers on the neighbour’s knollenorgans, which are dedicated to conspecific discharges and encode the emitter’s identity and bearing. This is the only one of the three that could be a *signal* in the strict sense.

Silencing a fish removes all three at once. Any behavioural consequence is therefore uninterpretable: it may be the muted fish foraging worse because it can no longer see, its neighbours foraging worse because they lost a source of free illumination, or its neighbours behaving differently because a signal they were attending to has disappeared. This confound is not specific to the model; it is the reason EOD-ablation experiments in real fish are hard to interpret.

Contributions. We make the three consequences independently switchable and ask what changes. (i) A **channel algebra** for dual-use signals: refference, illumination and detection become three masks on one motor act, letting us build pulses that are private, exploitable-but-undetectable, detectable-but-uninformative, phantom, replayed from another context, or scrambled in time or sender identity. (ii) An **interventional** causal influence: because we own the simulator, the counterfactual of Jaques et al. [8] is evaluated with the *do* operator directly—re-rendering the receiver’s input under $\text{do}(e_j=1)$ and $\text{do}(e_j=0)$ with every other cause on its factual path—rather than approximated with a model of other agents. (iii) A **cue/signal discriminator**: we train matched populations whose only difference is whether the focal agent can be *heard*, holding its own reception fixed, and measure the divergence of their emission policies against seed-to-seed divergence within each world. (iv) An **economics**: metabolic cost per discharge, predators that localise prey by discharges, cooperative versus competitive harvests, tunable social range, persistent identity. (v) An **assay validation**: dyads that are communicating by construction, against which every metric is checked before it is trusted, plus a **decoupled channel** and a **non-fish dual-use benchmark** to separate what is about electric fish from what is about dual use.

A note on wording. Our populations adapt within a training run by policy-gradient learning under a fixed reward; we do not run an evolutionary outer loop over heritable traits. We therefore write *emerge under learning* throughout, and reserve claims about evolutionary history for the Discussion, where we say explicitly what our design can and cannot license.

2 Related work

Measuring emergent communication. Most work equips agents with a dedicated, costless channel that has no other use, so that any use of it is communicative by construction and the open question is what code appears [10–13]. Within that setting, Lowe et al. [7] showed that message–behaviour association is insufficient evidence and introduced positive signalling and positive listening; Jaques et al. [8] proposed causal influence, estimated through a learned model of other agents; Eccles et al. [14] added training biases that make communication more likely to appear; Bogin et al. [15], Chaabouni et al. [16] study compositionality and context-independence. Our contribution to this line is not another metric but an audit: we show that these diagnostics are individually foolable, exhibit a case where positive signalling is *higher* in a population whose receivers are deaf and where causal influence *rises* as the channel becomes less informative, and validate the whole battery against dyads of known communication status before applying it.

Dual-use channels, and the cue/signal distinction. Our setting inverts the usual one: the channel is a physical act with a large private benefit, so the question is whether *any* of its use is communicative—the situation for active sonar, a robot’s laser sweep, or any action that both senses and is observable, and closer to the biological norm, since signals are generally thought to be ritualised from behaviours that had other functions. Noukhovitch et al. [17] show that influence under competition can reflect manipulation or cue-reading, the same identification problem we attack with the payoff and production tests. The distinction we operationalise is Scott-Phillips’s: a signal’s production has been shaped by its effect on receivers, a cue’s has not [4, 5]. Costly-signalling theory [18–20] explains how honesty is maintained in a signal that exists; our cost sweep asks whether cost brings one into being, and finds it does not.

Electric fish. Singh et al. [6] built the reference simulator and training setup we reimplement, reporting that discharges shape group foraging and spacing; our departure is definitional and causal, since their definition is effect-based and their one intervention removes three consequences at once. The biology we add is established: collective electrosensing [9], corollary-discharge cancellation [21], discharge energetics [22, 23], and electroreceptive eavesdroppers [24, 25].

3 Results

3.1 Do the assays work? A dyad that is communicating by construction

Every negative result below rests on an assay reporting “no communication here”. An assay that reported that unconditionally would produce the same paper, so before applying them we check that they fire on a case where communication is not in question.

We embed a Lewis signalling game in the same physics. One fish is immobile and privately observes which of two sites holds food; the other must travel to a site. The sender’s discharges are forced onto a fixed schedule, so pulse timing carries nothing, and the only free channel is a one-bit discharge subtype. Both fish are scored on the receiver reaching the correct site, and the harvest is shared. Crossing an informative sender (subtype = the private cue) against an uninformative one (a bit that is constant within a trial but uncorrelated with the truth), and a receiver that attends to the subtype against one that ignores it, gives four dyads whose communication status we know (Fig. 1). Only the informative-and-attended dyad succeeds: 4.00 correct arrivals out of four per episode, against 1.96 for chance.

Each assay behaves as it should, and—more usefully—each one alone is foolable. Positive signalling is 0.932 bits for the informative sender and 0.001 for the uninformative one, but it is just as high when nobody is listening, because it measures the sender in isolation. Causal influence and positive listening separate the attending receiver from the deaf one (31.5 versus 0.000 nats) but are just as high when the channel carries nothing, because a receiver can organise its behaviour around a variable that is pure noise. Even the payoff test is foolable on its own: deleting the subtype costs the attending receiver -172.5 when the subtype is informative and still -72.8 when it is noise, because the receiver has come to depend on it either way.

Only the conjunction identifies communication. This is the standard we hold the electric discharge to in what follows, and it is a standard the discharge could in principle meet: the instruments are demonstrably capable of detecting a payoff-relevant, informative, attended channel in this simulator.




Figure 1: The assays detect communication when it is present, and each is foolable alone. Four dyads in the referential game. Only *info+attend* is communicating; the other three each satisfy some subset of the diagnostics.

3.2 What the optimiser can and cannot do

A negative result about communication invites the objection that the learner, not the environment, is the limiting factor. We measured this rather than argued about it, because the answer bounds what the rest of the paper may claim.

In the same referential game we trained recurrent MAPPO under three conditions, twelve seeds each at 50×10^6 agent-steps: both sides learning; the receiver learning against a *scripted honest* sender, so the message is already perfectly informative and only reading it must be discovered; and both sides learning with the positive-signalling bias of Eccles et al. [14], which exists to break exactly this deadlock. *None of the thirty-six runs solved the game*: all sit at chance, against 4.00/4 for the hand-coded dyad reading the same observation vector. Reducing the steering noise by an order of magnitude and halving the credit horizon do not move it either (Appendix D).

So shared-parameter recurrent MAPPO—the standard configuration, and the one the reference model uses—does not discover that a single perfectly informative input bit should select between two known targets in an embodied task, even when told what to say. Our results therefore divide in two. The muting decomposition, the foolability of each diagnostic, and the metrics’ behaviour under intervention are measurements on trained policies and do not depend on this. The claims that metabolic cost fails to create a signal, and that a decoupled channel does not either, are claims about what learning arrives at and inherit the bound: these pressures did not produce signalling under a standard learner, not that they could not under a better one.

3.3 A validated, fast reimplementaion

The environment is a GPU reimplementaion of the reference simulator, checked stage by stage against the original NumPy code in double precision on random scenes. All 11 checks agree to floating-point round-off (maximum relative deviation 6×10^{-10} ; the knollenorgan and ampullary pathways agree bit-for-bit), and it sustains 596k agent-steps per second in aggregate on two GPUs, which is what makes 354 independent training runs affordable. Appendix A gives the per-stage table and records an incidental bug in the reference code: its configuration file specifies a wall-reflection coefficient of 0.95, but the scene builder never forwards it, so the published results in fact use unattenuated images.

3.4 The discharge is worth a great deal — to the emitter

Agents that can discharge eat 5.01 food items per episode; agents muted from the start of training eat 1.78 ($2.8\times$, Hedges’ $g = 4.3$, $p = 0.0001$). The probe does real work, which is the precondition for everything that follows: an EOD is not a cheap-talk token but an expensive and useful act of sensing that happens to be public. Removing the three consequences one at a time during training orders them—no detection 4.71, no illumination 4.31, private probe only 3.92, no refference 4.27—but from-scratch comparisons confound the value of a channel with the difficulty of learning without it, and training here is bimodal (Appendix B). We therefore rely on interventions applied to frozen policies.




Figure 2: Silencing one fish removes three things at once; separating them reassigns the effect. Within-arena paired differences from the intact channel, frozen policies. Significance is a permutation test over seeds.

3.5 The muting effect is almost entirely private sensing

The central experiment. We take a trained, intact population and, at evaluation time only, silence one fish—the intervention used upstream—then repeat with each of its three consequences removed separately. Every condition is run from the same episode seed, so arena k has the same initial food layout and the same sensory-noise stream in all of them, and differences are taken within arena.

Muting one fish costs that fish -0.52 food items per episode (Fig. 2A). Decomposed: removing only its own reafference costs -0.51, removing only its illumination of neighbours costs -0.00, and removing only its detectability costs 0.03. The private share accounts for essentially the whole effect.

The effect on its *neighbours* is smaller still (Fig. 2B): -0.01 items in total, with the illumination share 0.03 ($p = 0.5220$) and the detection share 0.01 ($p = 0.7234$). Group spacing (Fig. 2C) is likewise unchanged within resolution.

This is the paper’s first substantive result, and it is a negative one. When a fish is silenced in this model, what happens is overwhelmingly that the fish has been blinded, not that a conversation has been interrupted. An experiment that reports only the joint effect—as the reference study does, and as EOD-ablation experiments in real fish must—will attribute to communication an effect that is almost entirely private electrolocation.

3.6 The channel nonetheless carries information and moves policies

It would be wrong to conclude that the public channel is inert. It is not.

Discharge probability depends on the emitter’s situation: 0.586 when food is within 10 cm versus 0.457 when it is not ($p = 0.0001$; Fig. 3C). That state-dependence is inherited by the public channel, and a receiver can read it: decoding *only* from the number of pulses a receiver heard from a given emitter in a trailing window, food proximity is recovered with $\Delta R^2 = 0.156$ above a time-shifted baseline ($p = 0.0000$; Fig. 3B). Positive signalling about food, measured as mutual information conditioned on distance-to-neighbour decile so that geometry cannot manufacture it, is 0.086 of the state entropy (Fig. 3A).

Positive listening holds too. The interventional causal influence of Eq. 6—re-rendering a receiver’s sensory input under $\text{do}(e_j=1)$ versus $\text{do}(e_j=0)$ with everything else held on its factual path—is 0.0026 nats, against a noise floor of 0.00019 nats obtained by re-rendering the *same* intervention under an independent noise draw (Fig. 4A). The channel moves receiver policies 13× more than sensory noise does.

Both of these diagnostics, however, behave badly here in ways worth recording. Positive signalling is *higher* in populations whose receivers are deaf (0.152) than in populations whose receivers can hear (0.086): a fish modulates its probe rate with its own situation whether or not anybody is listening, so the measure is picking up the sensory schedule, not signalling. This is exactly the artifact Lowe et al. [7] warn of, and it is straightforward to demonstrate once a deaf population is available as a control. Causal influence misbehaves in the other direction: it *rises* 12-fold, from 0.0026 to 0.0303 nats, as metabolic cost is increased—that is, precisely as the pulse train becomes rarer and less informative (Fig. 4D). A rare discharge is more surprising,




Figure 3: What the pulse train carries. (A) Positive signalling, conditioned on distance-to-neighbour decile. (B) Cross-validated decoding from receiver-observable pulse counts only, minus the time-shifted baseline. (C) Discharge probability by food proximity.




Figure 4: Positive listening, and why it is not enough. (A) Interventional causal influence against its noise floor (log scale). (B) Positive listening against a rate-matched time-shift null and against silence, which is out of distribution. (C) Phantom-pulse dose-response on neighbours' intake. (D) The dissociation: as metabolic cost rises, causal influence per discharge *increases* while the information a receiver can decode *falls*. Influence and informativeness move in opposite directions, so neither substitutes for the other.

so switching it on or off moves the receiver's policy more, even though the train as a whole now tells the receiver less. Influence per message and information per unit time are different quantities, and in this system they are anti-correlated.

3.7 ...but it has no payoff consequence

Positive signalling and positive listening are both satisfied, so on the effect-based definition this is communication. The payoff test says otherwise.

Replacing an emitter's public train with the train the same policy produced in a *different arena*—identical marginal rate, burst structure and inter-pulse statistics, no contingency with the receiver's actual world—changes neither party's return. The same holds for temporal scrambling, for identity scrambling, and for phantom pulses injected at rates from 0.05 to 0.6 (Fig. 4C). Because a non-significant difference is not evidence of no difference, we report these as two one-sided tests against an equivalence margin fixed in advance at 5% of the intact condition's own return: no comparison clears the margin in either direction — the intervals are simply wider than the ± 2.13 margin (cross-arena replay, $+0.77$ (90% CI $[-2.53, +4.09]$); temporal scrambling, $+0.35$ (90% CI $[-2.96, +3.68]$); removing the public channel entirely, -0.75 (90% CI $[-4.11, +2.57]$)). We therefore report these as *no effect detectable at this sample size*, and explicitly not as demonstrated equivalence. The honest statement is that the design rules out payoff effects larger than roughly a tenth of the intact return, and is not powered to rule out smaller ones.

The contingency of the public channel, as opposed to its presence, is worth approximately nothing to anybody. A receiver can read the emitter's situation off its pulse train, and does adjust its policy when the




Figure 5: Sender shaping measured with reception held fixed. (A) SSI from the yoked contrast, which cuts audibility only, beside the confounded deaf contrast, which cuts reception. (B) Discharge rate, live versus yoked. (C) The free-riding effect that drives the deaf contrast: fish that can hear neighbours discharge less.

train changes, but nothing downstream of that adjustment shows up in what either fish eats.

3.8 Cost and predation regulate the probe without creating a signal

The reference model gives discharges no cost. Adding one produces a clean dose-response: discharge probability falls monotonically from 0.495 at $\lambda = 0$ to 0.032 at $\lambda = 0.08$ (Fig. 8). Eavesdropping predation pushes it lower again at every cost, and mortality falls with it, so crypsis works. Cooperative harvesting sustains a higher rate than competition at matched cost.

We expected this to make the pulse train informative. A budget has to be spent where it pays, and a schedule that allocates pulses to the moments that matter should be more diagnostic of the emitter’s state than a constant one. The opposite happens. Information about food falls monotonically with cost, from $\Delta R^2 = 0.098$ at $\lambda = 0$ to 0.010 at $\lambda = 0.08$, and the same decline appears for every other target we decoded. It survives normalising per pulse: informativeness per discharge falls from 0.127 to 0.087. The learned response to a metabolic price is a lower baseline rate, not a smarter schedule.

Silence behaves the same way. Discharge probability is essentially identical whether or not a predator is within 25 cm (difference -0.004 , 95% CI $[-0.007, +0.000]$), and danger is close to undecodable from the public channel ($+0.010$ AUC above the time-shifted baseline). Agents manage their average exposure rather than timing their gaps, which is what a prey animal that cannot see the predator coming can do; our predator detects discharges at 60 cm but is itself passively detectable only at about 25 cm.

3.9 Sender shaping, with reception held fixed

Our first attempt at the production-side test compared populations trained with and without knollenorgans, and gave SSI = 0.415 at zero cost. That contrast is confounded, and in a way worth spelling out because it is easy to make. Disabling knollenorgans removes *reception*. It therefore asks whether a fish behaves differently when it can hear others, not whether it behaves differently when others can hear *it*. The direction of the effect gives the confound away: hearing populations discharge *less* than deaf ones (0.545 versus 0.611), which is free-riding—a fish that can hear its neighbours’ probes cuts its own costly sampling—and not audience design, which would push emission up or restructure it.

The yoked condition holds reception fixed and cuts audibility instead. Receivers get knollen input with matched statistics, drawn from another arena, so no sender’s pulses reach anybody while every receiver still has a realistic social input stream. Measured this way, sender shaping is 0.042 at zero cost (Fig. 5A), and discharge rate is 0.576 live versus 0.583 yoked. Emission policy is therefore no more different between the live and yoked worlds than between seeds of the same world. Holding reception fixed, being listened to leaves no detectable trace on how the discharge is scheduled, which is what it means for the pulse to be a cue rather than a signal.

Table 1: The muting decomposition across scales and architectures. “Private share” is the fraction of the total muting effect attributable to the emitter’s own lost refference.

Condition	seeds	total	−refference	−detection	private share
baseline (4 fish, 60 cm, 512 steps)	12	-0.52	-0.51	+0.03	98%
2 fish	6	-1.66	-1.58	-0.02	95%
6 fish	6	-0.43	-0.59	-0.08	135%
100 cm arena	6	-0.31	-0.34	-0.02	107%
1024-step episodes	6	-1.10	-0.92	+0.08	83%
256-unit GRU	6	-0.61	-0.57	+0.11	93%

3.10 A decoupled variable is what lets a social function appear

A negative result is only as good as the positive control beside it, and to this point every manipulation has left the discharge welded to its sensory function. We therefore gave the agents a second variable: a discharge subtype, one bit, that rides on the pulse—so sending still costs a pulse and still exposes the fish to predators—but whose value changes nothing about the emitter’s own electric image, nothing about how the pulse illuminates neighbours, and nothing about detectability. Only the content is decoupled, not the act.

Agents mark 52% of their discharges with the subtype, which is indistinguishable from the 50% an unmodulated bit would give, and it carries 0.0141 bits about food proximity above its permutation null. Deleting the subtype leaves the pulse, and therefore every sensory consequence, completely intact, so any payoff change is attributable to content alone. Across the four ecological settings: standard $\times$ compet., +0.51 (95% CI $[-0.45, +1.40]$, $p = 0.292$); standard $\times$ cooper., +0.54 (95% CI $[-1.81, +2.99]$, $p = 0.640$); sparse $\times$ compet., −0.76 (95% CI $[-1.76, +0.20]$, $p = 0.165$); sparse $\times$ cooper., +0.51 (95% CI $[-0.88, +1.95]$, $p = 0.512$). No cell survives correction for multiple comparisons. Freeing the content of a discharge from its sensory function was not, on its own, sufficient to make that content payoff-relevant under any pressure we applied.

3.11 The same decomposition in a world with no electric fish in it

If the result is about dual-use action channels rather than about electroreception, it should reappear wherever an action senses and broadcasts at once. We built a minimal environment with no fish physics in it: agents forage in a unit square, and a PING action has exactly the three consequences a discharge has—it shows the pinger the items near itself, it shows nearby others the items near *them*, and it is registered by everyone as an event with a bearing. A weak always-on passive sense means silencing an agent is not the same as blinding it. The same masks, the same trainer and the same intervention battery apply unchanged.

The decomposition replicates. Agents that can ping collect 13.07 items per episode against 6.03 for permanently silent ones; muting one agent costs it −9.89 items, of which −9.89 (100%) is its own lost refference and −0.06 its lost detectability (Appendix F). Pricing the ping suppresses it from 0.753 to 0.001. Neither the sign nor the rough magnitude of the decomposition depends on the electric-fish biophysics; what it depends on is that one action does two jobs.

3.12 The regime map, and how far it generalises

Appendix C gives the full cost $\times$ predation $\times$ economics grid and the rule-based regime labels. The structure is orderly, but it is the structure of a probe schedule being turned up and down: the axes that move are discharge rate and mortality, while the informational and payoff-relevance axes stay near zero throughout.

The decomposition is not an artefact of one configuration. Repeating it with two and six fish, a 100 cm arena, 1024-step episodes and a 256-unit recurrent policy, the private share of the muting effect stays between 83% and 135% (Table 1), and the detection share remains indistinguishable from zero in every case.

4 Discussion

An effect on a receiver is not a signal. The operational definition used by Singh et al. [6]—an EOD that conspecifics detect and that changes their behaviour—is satisfied in our model at every metabolic cost we tested, including zero. It is satisfied because the knollenorgan channel really does carry information and receivers really do respond to it. Yet the same channel has no detectable consequence for anyone’s payoff: replaying, scrambling or fabricating it moves neither party’s return by an amount we can resolve. We are careful about what that licenses. Our equivalence tests do not clear the 5% margin we set in advance, because the confidence intervals are wider than the margin; the defensible statement is *no effect detectable at this sample size*, not demonstrated equivalence. What makes the negative interpretable is not the width of those intervals but the assay validation of §3.1: the same tests, applied to a dyad that is communicating by construction, return effects one to two orders of magnitude larger than anything we see here.

So the strongest defensible form of the claim is this. A system can satisfy the standard positive-signalling and positive-listening diagnostics without any detectable payoff relevance and without evidence that production was shaped by being heard.

Establishing the production-side half of the argument turned out to need more care than we first gave it. Disabling knollenorgans removes *reception*, so a hearing/deaf contrast asks whether a fish behaves differently when it can hear others—not whether it behaves differently when others can hear *it*. Those are different questions, and the first is contaminated by free-riding: a fish that can hear its neighbours’ probes cuts its own costly sampling, which is what we in fact observe. The control that isolates being listened to has to hold reception fixed and cut audibility, which is what our yoked condition does.

The muting confound. Silencing a fish is the natural intervention and it is the one available in real experiments, but it removes three things simultaneously. Our decomposition shows that the three shares are not merely different in size—they act on different individuals and in different directions. Losing refference hurts the muted fish; losing the illumination it provided hurts its neighbours; losing its detectability changes how the group is spaced. A single number attributed to “communication” averages over all three. The same caution applies to real ablation work: an EOD-silenced *Gnathonemus* is not a fish with its telephone taken away, it is a fish that has been blinded and has stopped lighting the room for everyone else.

Cost does not manufacture semantics. It is natural to expect that pricing a pulse would turn a probe schedule into an informative one: a budget has to be spent where it pays, and that allocation should make the train diagnostic of the emitter’s state. That is not what happens. Across the whole cost range the pulse train becomes rarer and *less* informative about every target we decoded, and it stays less informative after normalising per pulse. The learned response to a price is a lower baseline rate, not a smarter schedule. Whatever costly signalling does in the classical models [18–20], it presupposes a signal whose honesty is at stake; it does not by itself bring one into being.

What we should not conclude is that a binary pulse leaves no room for a code. A binary event repeated in continuous time supports a large space of temporal codes—inter-pulse intervals, bursts, pauses, event-locked patterns—and mormyrid communication is known to use pulse timing and waveform identity [3, 26, 27]. Our agents *failed to exploit* the temporal code available to them under the pressures we applied. That is a claim about what learning found, not about what the channel permits.

Silence tracks the regime, not the moment. Cessation of discharge is documented in mormyrids as both predator avoidance and submissive display [24, 28], and we expected a momentary version: quiet when the predator is near. It does not appear. Discharge probability is essentially unchanged by a predator within 25 cm (−0.004, 95% CI [−0.007, +0.000]) and danger is near-undecodable from the public channel (+0.010 AUC). Agents lower their *baseline* rate wherever predators exist, and mortality falls with it, so crypsis works without timing. The likely reason is sensory asymmetry: our predator hears discharges at 60 cm but is passively detectable only at 25 cm.

Signalling as a public good and a public hazard. Two externalities operate in opposite directions, neither in the reward function: a discharge illuminates neighbours’ worlds, and it also jams their passive sense and draws an eavesdropping predator toward the group. Cooperative harvesting sustains a higher rate than

competition at matched cost, consistent with the public-good component being internalised when the harvest is shared. We do not state that as a mechanism: our interventions found no detectable payoff consequence of the public channel, so we cannot also claim neighbours benefit enough from a pulse to explain the rate difference, which may instead reflect altered encounter distributions or an effect below our resolution.

What the discharge lacks is not a channel. The decoupled control was meant to test the hypothesis that dual use is what blocks signalling: give the fish a variable it can move without degrading its own sensing, and a social function should appear. It does not. Agents mark half their discharges with the subtype—the rate an unmodulated bit would give—and deleting it costs nobody anything in any of the four ecological settings.

Setting that beside the referential assay says what the missing ingredient is. The assay differs from the foraging world in three ways, and only one of them is the channel. It also gives the sender information the receiver cannot get for itself, and it pays both of them when the receiver acts on it. In the foraging world a fish can find food with its own probe, so a neighbour’s private state is worth very little; there is a channel available and nothing to say over it. Channel factorisation is plausibly necessary for a signal to become distinct from a probe, but it is clearly not sufficient. That is a testable prediction for the biology: electric communication should be elaborated in exactly those contexts where one fish reliably knows something a conspecific cannot discover first-hand—dominance status, reproductive state, individual identity—and not in contexts where both parties can simply sense the world.

Relation to emergent-communication work. Most emergent-communication studies equip agents with a dedicated, costless, otherwise-useless channel [10–12]. Under those conditions any use of the channel is communicative by construction, and the interesting question is what code emerges. Electric fish invert the setup: the channel is a physical act with a large private benefit, and the interesting question is whether *any* of its use is communicative. We think this is the more common situation in biology—most signals are thought to be ritualised from behaviours that originally had other functions—and it is a setting in which the standard metrics need the production-side test that SSI supplies. Positive signalling and positive listening are both necessary; neither, nor both together, is sufficient.

5 Limitations

Six bound the claims above and are set out in full in Appendix G. The three that matter most: our populations adapt by policy-gradient learning within a training run, not by an evolutionary outer loop, so the Sender Shaping Index is a within-lifetime analogue of Scott-Phillips’s criterion rather than the criterion itself, and we write *emerge under learning* throughout. Our knollenorgan channel is all-or-none, as upstream, so graded amplitude signalling is structurally impossible in the coupled condition, and the decoupled condition lifts only the crudest version of that restriction — one bit, not a waveform. And the equivalence tests do not reach the margin we set in advance; what makes the nulls interpretable is the assay validation of §3.1, not the width of those intervals.

References

- [1] H. W. Lissmann and K. E. Machin. Electric receptors in a non-electric fish (Clarias). *Nature*, 199(4888):88–89, 1963. doi: 10.1038/199088a0.
- [2] Gerhard von der Emde. Active electrolocation of objects in weakly electric fish. *Journal of Experimental Biology*, 202(10):1205–1215, 1999. doi: 10.1242/jeb.202.10.1205.
- [3] Carl D. Hopkins. Neuroethology of electric communication. *Annual Review of Neuroscience*, 11(1):497–535, 1988. doi: 10.1146/annurev.ne.11.030188.002433.
- [4] John Maynard Smith and David Harper. *Animal Signals*. Oxford Series in Ecology and Evolution. Oxford University Press, Oxford, 2003. ISBN 9780198526858.

- [5] Thomas C. Scott-Phillips. Defining biological communication. *Journal of Evolutionary Biology*, 21(2):387–395, 2008. doi: 10.1111/j.1420-9101.2007.01497.x.
- [6] Satpreet H. Singh, Sonja Johnson-Yu, Zhouyang Lu, Aaron Walsman, Federico Pedraja, Denis Turcu, Pratyusha Sharma, Naomi Saphra, Nathaniel B. Sawtell, and Kanaka Rajan. Active electrosensing and communication in MARL-trained weakly electric fish collectives, 2025.
- [7] Ryan Lowe, Jakob Foerster, Y-Lan Boureau, Joelle Pineau, and Yann Dauphin. On the pitfalls of measuring emergent communication. In *Proceedings of the 18th International Conference on Autonomous Agents and MultiAgent Systems (AAMAS)*, pages 693–701. International Foundation for Autonomous Agents and Multiagent Systems, 2019.
- [8] Natasha Jaques, Angeliki Lazaridou, Edward Hughes, Caglar Gulcehre, Pedro A. Ortega, DJ Strouse, Joel Z. Leibo, and Nando de Freitas. Social influence as intrinsic motivation for multi-agent deep reinforcement learning. In *Proceedings of the 36th International Conference on Machine Learning (ICML)*, volume 97 of *Proceedings of Machine Learning Research*, pages 3040–3049, 2019.
- [9] Federico Pedraja and Nathaniel B. Sawtell. Collective sensing in electric fish. *Nature*, 628(8006):139–144, 2024. doi: 10.1038/s41586-024-07157-x.
- [10] Jakob N. Foerster, Yannis M. Assael, Nando de Freitas, and Shimon Whiteson. Learning to communicate with deep multi-agent reinforcement learning. In *Advances in Neural Information Processing Systems 29 (NIPS)*, pages 2137–2145, 2016.
- [11] Igor Mordatch and Pieter Abbeel. Emergence of grounded compositional language in multi-agent populations. In *Proceedings of the Thirty-Second AAAI Conference on Artificial Intelligence (AAAI-18)*, pages 1495–1502, 2018. doi: 10.1609/aaai.v32i1.11492.
- [12] Angeliki Lazaridou, Alexander Peysakhovich, and Marco Baroni. Multi-agent cooperation and the emergence of (natural) language. In *International Conference on Learning Representations (ICLR)*, 2017.
- [13] Angeliki Lazaridou, Karl Moritz Hermann, Karl Tuyls, and Stephen Clark. Emergence of linguistic communication from referential games with symbolic and pixel input. In *International Conference on Learning Representations (ICLR)*, 2018.
- [14] Tom Eccles, Yoram Bachrach, Guy Lever, Angeliki Lazaridou, and Thore Graepel. Biases for emergent communication in multi-agent reinforcement learning. In *Advances in Neural Information Processing Systems 32 (NeurIPS)*, pages 13111–13121, 2019.
- [15] Ben Bogin, Mor Geva, and Jonathan Berant. Emergence of communication in an interactive world with consistent speakers, 2018. Emergent Communication Workshop, NeurIPS 2018.
- [16] Rahma Chaabouni, Eugene Kharitonov, Diane Bouchacourt, Emmanuel Dupoux, and Marco Baroni. Compositionality and generalization in emergent languages. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics (ACL)*, pages 4427–4442, 2020. doi: 10.18653/v1/2020.acl-main.407.
- [17] Michael Noukhovitch, Travis LaCroix, Angeliki Lazaridou, and Aaron Courville. Emergent communication under competition. In *Proceedings of the 20th International Conference on Autonomous Agents and MultiAgent Systems (AAMAS)*, pages 974–982. International Foundation for Autonomous Agents and Multiagent Systems, 2021.
- [18] Amotz Zahavi. Mate selection—a selection for a handicap. *Journal of Theoretical Biology*, 53(1):205–214, 1975. doi: 10.1016/0022-5193(75)90111-3.
- [19] Alan Grafen. Biological signals as handicaps. *Journal of Theoretical Biology*, 144(4):517–546, 1990. doi: 10.1016/S0022-5193(05)80088-8.
- [20] Szabolcs Számadó. The cost of honesty and the fallacy of the handicap principle. *Animal Behaviour*, 81(1):3–10, 2011. doi: 10.1016/j.anbehav.2010.08.022.
- [21] Avner Wallach and Nathaniel B. Sawtell. An internal model for canceling self-generated sensory input in freely behaving electric fish. *Neuron*, 111(16):2570–2582.e5, 2023. doi: 10.1016/j.neuron.2023.05.019.
- [22] Vielka L. Salazar, Rüdiger Krahe, and John E. Lewis. The energetics of electric organ discharge generation in gymnotiform weakly electric fish. *Journal of Experimental Biology*, 216(13):2459–2468, 2013. doi: 10.1242/jeb.082735.

- [23] Michael R. Markham, M. Lynne McAnelly, Philip K. Stoddard, and Harold H. Zakon. Circadian and social cues regulate ion channel trafficking. *PLoS Biology*, 7(9):e1000203, 2009. doi: 10.1371/journal.pbio.1000203.
- [24] Philip K. Stoddard. Predation enhances complexity in the evolution of electric fish signals. *Nature*, 400(6741): 254–256, 1999. doi: 10.1038/22301.
- [25] S. Hanika and Bernd Kramer. Electric organ discharges of mormyrid fish as a possible cue for predatory catfish. *Naturwissenschaften*, 86(6):286–288, 1999. doi: 10.1007/s001140050616.
- [26] Christa A. Baker, Tsunehiko Kohashi, Ariel M. Lyons-Warren, Xiaofeng Ma, and Bruce A. Carlson. Multiplexed temporal coding of electric communication signals in mormyrid fishes. *Journal of Experimental Biology*, 216(13): 2365–2379, 2013. doi: 10.1242/jeb.082289.
- [27] Bruce A. Carlson and Carl D. Hopkins. Stereotyped temporal patterns in electrical communication. *Animal Behaviour*, 68(4):867–878, 2004. doi: 10.1016/j.anbehav.2003.10.031.
- [28] Peter Moller, Jacques Serrier, and Deborah Bowling. Electric organ discharge displays during social encounter in the weakly electric fish *Brienomyrus niger* L. (Mormyridae). *Ethology*, 82(3):177–191, 1989. doi: 10.1111/j.1439-0310.1989.tb00498.x.
- [29] Sat Gavassa, Ana C. Silva, Emmanuel Gonzalez, and Philip K. Stoddard. Signal modulation as a mechanism for handicap disposal. *Animal Behaviour*, 83(4):935–944, 2012. doi: 10.1016/j.anbehav.2012.01.012.
- [30] Stefan Schuster. Count and spark? The echo response of the weakly electric fish *Gnathonemus petersii* to series of pulses. *Journal of Experimental Biology*, 204(8):1401–1412, 2001. doi: 10.1242/jeb.204.8.1401.
- [31] Ling Chen, Jonathan L. House, Rüdiger Krahe, and Mark E. Nelson. Modeling signal and background components of electrosensory scenes. *Journal of Comparative Physiology A*, 191(4):331–345, 2005. doi: 10.1007/s00359-004-0587-3.
- [32] Chao Yu, Akash Velu, Eugene Vinitzky, Jiakuan Gao, Yu Wang, Alexandre Bayen, and Yi Wu. The surprising effectiveness of PPO in cooperative multi-agent games. In *Advances in Neural Information Processing Systems 35 (NeurIPS) Datasets and Benchmarks Track*, volume 35, pages 24611–24624, 2022.
- [33] John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms, 2017.
- [34] Kyunghyun Cho, Bart van Merriënboer, Caglar Gulcehre, Dzmitry Bahdanau, Fethi Bougares, Holger Schwenk, and Yoshua Bengio. Learning phrase representations using RNN encoder–decoder for statistical machine translation. In *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 1724–1734, 2014. doi: 10.3115/v1/D14-1179.
- [35] John Schulman, Philipp Moritz, Sergey Levine, Michael I. Jordan, and Pieter Abbeel. High-dimensional continuous control using generalized advantage estimation. In *International Conference on Learning Representations (ICLR)*, 2016.
- [36] Richard Dawkins and John R. Krebs. Animal signals: information or manipulation? In John R. Krebs and Nicholas B. Davies, editors, *Behavioural Ecology: An Evolutionary Approach*, pages 282–309. Blackwell Scientific Publications, Oxford, 1 edition, 1978.

Table 2: Agreement between the GPU reimplementation and the reference NumPy implementation, in double precision on random scenes. Error is the maximum absolute deviation normalised by the largest magnitude present.

Stage	Max. relative deviation
Coulomb field of point charges	2.3×10^{-15}
Field of point dipoles	1.4×10^{-14}
Field including first-order wall images	1.9×10^{-15}
Dipoles induced on conductors	1.2×10^{-15}
Charge-limiting of fish body moments	1.2×10^{-15}
Corollary-discharge baseline	1.6×10^{-14}
Ampullary intrinsic baseline	1.3×10^{-14}
Mormyromast, self-image mode	6.0×10^{-10}
Mormyromast, collective-sensing mode	4.9×10^{-11}
Knollenorgan (conspicific detection)	exact (0)
Ampullary (passive DC)	exact (0)

Appendix

A Numerical validation against the reference implementation

Every stage of the sensory pipeline is checked against the original NumPy code in double precision on random scenes (`scripts/validate_physics.py`). Agreement is at floating-point round-off, and the knollenorgan and ampullary pathways agree bit-for-bit.




Figure 6: Maximum relative deviation from the reference NumPy implementation over random scenes, in double precision. Bars at the left edge are exactly zero.

An incidental bug in the reference code. Its configuration file sets a wall-reflection coefficient of 0.95, but `ElectricScene._build_reflections` never forwards it to `reflect_sources`, whose default is 1.0. The published results therefore use unattenuated first-order images. We match the code rather than the configuration file; using 0.95 shifts the corollary-discharge baseline by $\sim 1.4 \times 10^{-6}$ relative and changes ampullary readings qualitatively, because those are a small difference of large numbers.

B Training the three channel conditions from scratch

Training is bimodal: a seed either discovers patch foraging or stays near the floor set by passive sensing alone. A seed counts as converged if it reaches half the ceiling of its own task family, a criterion applied identically to every condition within a family; the fraction clearing it is reported as its own outcome, because conditions that remove sensory channels are not only worse on average but less reliably learnable.




Figure 7: Removing each consequence of a discharge, trained from scratch. Bars are means over independent seeds with bootstrap 95% CIs; points are seeds. Percentages in (A) give the fraction of seeds reaching the learning criterion.

C Regime classification




Figure 8: Metabolic cost and eavesdropping predation jointly set discharge policy. Rows are competitive and cooperative harvests; each cell is a mean over independent seeds.

D Why the learned referential game fails

The hand-coded receiver reads the subtype and the two site bearings out of the observation vector and reaches the correct site on every trial, so the information is present and a competent policy exists inside the representable class. To check that the learned failure is not an artefact of exploration noise or of the delay between committing to a direction and arriving, we re-ran the receiver-only condition (scripted honest sender) with the initial policy standard deviation on the steering action reduced from $e^{-0.5}$ to e^{-2} , and with the two sites moved from 60 cm apart at a range of 49 cm to 40 cm apart at a range of 25 cm, which roughly halves the credit-assignment horizon and doubles the number of trials per episode. Neither change moves




Figure 9: Rule-based regime labels from the pooled metric vector. Thresholds are percentiles of each metric’s own distribution across the grid rather than hand-set values, so the boundaries describe a continuum and are not claims about sharp transitions.




Figure 10: A hand-coded dyad solves the referential game from the same observations; twelve seeds each of three learned configurations, including one with a published bias designed for this problem, do not leave chance.

performance off chance:

Configuration	arrivals / episode	chance
far sites, low steering noise	2.01 of 4	2.0
near sites, default noise	3.00 of 6	3.0
near sites, low steering noise	3.07 of 6	3.0

We therefore attribute the failure to the difficulty of discovering a multiplicative interaction between a single binary input and a continuous steering policy under policy-gradient updates, rather than to the task geometry or the exploration schedule.

E The decoupled signalling variable

F The dual-use benchmark




Figure 11: A discharge subtype decoupled from sensing, crossed with competitive and cooperative harvests and with a sparse-patch task in which finding food is the bottleneck. (C) deletes the subtype while leaving the pulse, and therefore all sensing, completely intact — the manipulation the coupled channel does not admit.




Figure 12: A minimal dual-use foraging world with no fish physics: a PING senses, illuminates neighbours and is registered as an event. The muting effect decomposes the same way as in the electric-fish environment, and cost suppresses the ping the same way.

G Limitations in full

We list these plainly because several of them bound the claims above.

1. **Learning is not evolution.** Our populations adapt within a training run by policy-gradient learning under a fixed reward. The Scott-Phillips criterion refers to evolutionary history, and SSI is a within-lifetime analogue: it asks whether the emission policy that learning arrives at depends on whether receivers respond. A genuine evolutionary outer loop over heritable emission traits would test the criterion directly, and we did not run one. We write “emerge under learning” throughout rather than “evolve”.
2. **One arena scale, one group size.** Everything is measured with four fish in a 60×60 cm arena, in which the knollenorgan channel at unit gain reaches every conspecific. The range sweep partially addresses this, but group size is not varied, and the upstream model randomises arena size in a way we deliberately do not.
3. **Pulses are binary and fixed-amplitude.** Mormyrid electric communication uses inter-pulse-interval patterning and species- and individual-specific waveform structure [26, 27]; in wave-type gymnotiforms, amplitude modulations carry honest size information [29]. Our knollenorgan channel is all-or-none, as upstream, so waveform and amplitude signalling are structurally impossible in the coupled condition. This is precisely the limitation the decoupled condition is designed to lift, and it lifts only the crudest version of it: one bit of subtype, not a graded waveform.
4. **Episodes are short.** 512 steps is about 6s of fish time. Dominance relationships in real fish are established over much longer horizons, so the dominance-related results should be read as within-encounter,

not as stable social structure.

5. **The metabolic cost is a reward term, not a physiological model.** We express cost in reward units and report the fraction of the achieved return it consumes so that it can be compared with measured energy budgets, but we do not model the electrocyte physiology that sets the true cost, nor its dependence on amplitude.
6. **The predator is a caricature.** It is guided by discharges alone. Real electroreceptive predators also use vision, mechanoreception and olfaction, so real silence buys less safety than it does here, and the cryptic regime is correspondingly easier to reach in our model than in a river.
7. **Regime labels are a classification imposed on a continuum.** The phase diagram is generated by thresholding continuous metrics at percentiles of their own pooled distribution. The boundaries are therefore descriptive, not sharp transitions, and we make no claim about criticality or bifurcation.

H Methods

H.1 Simulator

The environment is a GPU reimplementaion of the weakly-electric-fish simulator of Singh et al. [6], rewritten so that thousands of independent arenas are advanced as a single set of tensor operations. All electrostatic constants, receptor counts, receptor placements, dynamic ranges, transduction functions, locomotor limits and reward coefficients are transcribed from that codebase; the differences are enumerated in §H.6.

State and kinematics. Each arena holds $F = 4$ fish, $N = 48$ food items and optionally one predator, in a $60 \text{ cm} \times 60 \text{ cm}$ rectangular tank with insulating walls. Fish i has position $\mathbf{x}_i$, heading θ_i and a body size $s_i \in [0, 1]$ that is fixed within an episode and sets a dominance rank. The timestep is $\Delta t = 1/83 \text{ s} = 12.05 \text{ ms}$, matched by the original authors to the social echo-response latency of *Gnathonemus petersii* [30]; episodes last 512 steps ($\approx 6.2 \text{ s}$). Motion is first-order in the commanded velocities,

$$\omega_i = u_i^{\text{turn}} \omega_{\max}(1 + s_i), \quad \mathbf{x}_i \leftarrow \mathbf{x}_i + \hat{u}(\theta_i) \max(u_i^{\text{fwd}}, 0) v_{\max}(1 + s_i), \quad (1)$$

with $v_{\max} = 0.422 \text{ cm per step}$ and $\omega_{\max} = 0.0434 \text{ rad per step}$, taken from the empirical percentiles used upstream. A fish whose proposed position would overlap another fish does not move and receives a collision penalty; walls clamp position without penalty.

Electric fields. All fields are exact Coulomb superposition, with the distance softening $\varepsilon_m = 10^{-5} \text{ m}$ applied *before* exponentiation exactly as upstream. A discharging fish carries two monopoles at $\pm 0.5 \text{ cm}$ along its body axis with charges $\pm q_0$, $q_0 = 1.11 \times 10^{-15} \text{ C}$ [31]:

$$\mathbf{E}_{\text{mono}}(\mathbf{r}) = k \sum_p \frac{q_p(\mathbf{r} - \mathbf{r}_p)}{(\|\mathbf{r} - \mathbf{r}_p\| + \varepsilon_m)^3}, \quad \mathbf{E}_{\text{dip}}(\mathbf{r}) = k \sum_d \left[\frac{3(\mathbf{p}_d \cdot \mathbf{u}_d)\mathbf{u}_d}{\|\mathbf{u}_d\|^5} - \frac{\mathbf{p}_d}{\|\mathbf{u}_d\|^3} \right], \quad (2)$$

with $k = 8.99 \times 10^9 \text{ N m}^2 \text{ C}^{-2}$ and $\mathbf{u}_d = \mathbf{r} - \mathbf{r}_d$. Fish and food carry permanent intrinsic dipoles (1.11×10^{-23} and $1.11 \times 10^{-24} \text{ C m}$). Food items and non-discharging fish are polarisable insulating spheres,

$$\mathbf{p}_n^{\text{ind}} = 3\varepsilon_0 V_n \chi \mathbf{E}(\mathbf{y}_n), \quad V_n = \frac{4}{3}\pi a_n^3, \quad \chi = -0.5, \quad (3)$$

with fish moments capped at $q_0 a_n$. Walls contribute one order of images across each of the four faces, unattenuated and without sign flip.

Receptors. Each fish carries 36 mormyromasts (10 concentrated in a $\pi/3$ chin fovea, 26 spread over the body), 24 ampullary organs and 12 knollenorgans *per conspecific*. A receptor reads the projection of the local field onto its outward normal and transduces it by a sign-preserving logarithmic compression into $[-1, 1]$,

$$\hat{s} = \text{sgn}(s) \frac{\log_{10} \text{clip}(|s|, E_{\min}, E_{\max}) - \log_{10} E_{\min}}{\log_{10} E_{\max} - \log_{10} E_{\min}}. \quad (4)$$

Mormyromasts use $[E_{\min}, E_{\max}] = [5 \times 10^{-8}, 5 \times 10^{-2}] \text{ V m}^{-1}$ and report the field of the *induced* dipoles alone: the direct discharge field is cancelled, so what remains is the electric image of objects. Images carried by a conspecific’s discharge rather than one’s own are rescaled by 100, as upstream. Ampullary organs ($[2 \times 10^{-10}, 2 \times 10^{-8}] \text{ V m}^{-1}$) sense only the permanent intrinsic dipoles and are therefore independent of any discharge; a conspecific discharge degrades them (multiplicative noise fraction rising from 0.05 to 0.5), so emitting jams a neighbour’s passive sense. Knollenorgans are all-or-none: a slot fires with the sign of the local field when $|s| > 2 \times 10^{-7} \text{ V m}^{-1}$ (a detection range of $\approx 100 \text{ cm}$ at unit gain). Crucially, the knollenorgan array is indexed *by emitter*, so sender identity and bearing are preserved, and each slot carries a noisy scalar of the emitter’s relative body size.

Task and reward. Food is distributed in four Gaussian patches ($\sigma = 6 \text{ cm}$) and is not replenished. A fish eats an item within 2 cm and $\pm\pi/4$ of its heading; contested items go to the larger fish. The per-step reward is

$$r_i = 10 n_i^{\text{eat}} + \frac{d_i^{\text{prev}} - d_i^{\text{food}}}{W+H} - 5 b_i^{\text{bitten}} (1 + s_{\text{biter}} - s_i) - 0.001 b_i^{\text{bite}} - 0.5 c_i^{\text{coll}} \underbrace{- \lambda e_i}_{\text{new}} \underbrace{- \rho k_i}_{\text{new}}, \quad (5)$$

where $e_i \in \{0, 1\}$ indicates a discharge and k_i a predator strike. The first five terms are upstream’s; λ (metabolic cost) and ρ (predation) are introduced here. Under cooperative harvesting a fraction α of each neighbour’s intake is added to r_i .

Predators. A predator perceives only fish that discharged on the current step and lie within 60 cm, weights them by $1/(d+1)$, and advances 0.45 cm toward the loudest. Within 4 cm it strikes the nearest fish—not necessarily the emitter—and enters a 40-step refractory period before respawning at a random location. A silent fish is invisible, but a silent fish beside a discharging neighbour is not safe, so signalling imposes a risk externality on the group. The predator carries a strong intrinsic dipole and is therefore detectable passively at $\approx 25 \text{ cm}$, far inside the range at which it detects discharges.

H.2 The channel algebra

A discharge by fish j at time t is routed through three independent binary masks:

V_j^{self}	does j ’s pulse produce j ’s own electric image?
V_j^{illum}	does j ’s pulse produce images on <i>other</i> fish’s mormyromasts?
V_j^{det}	does j ’s pulse register on other fish’s knollenorgans?

Setting all three equal to the emitted action recovers the biological default. The conditions used here are: PRIVATE ($V^{\text{illum}} = V^{\text{det}} = 0$), CUE-ONLY ($V^{\text{det}} = 0$), SIGNAL-ONLY ($V^{\text{illum}} = 0$), SOCIAL-ONLY ($V^{\text{self}} = 0$), and MUTE (all zero, the upstream intervention). Three further manipulations act on what receivers hear without changing what the emitter does, so that any change in receiver behaviour is attributable to the public channel alone:

- **Phantom.** Receivers perceive discharges from j drawn from an exogenous Bernoulli process that j did not choose.
- **Replay.** j ’s public train is replaced by the train the same policy produced in a *different arena*. Marginal rate, burst structure and inter-pulse-interval statistics are preserved exactly—these are the same trains—while the contingency between a pulse and the receiver’s current world is destroyed. A context-swapped variant pairs each arena with the one whose food-proximity time series is most opposite.
- **Scrambling.** Timing scrambling circularly shifts each train, preserving rate and destroying temporal alignment; identity scrambling permutes the knollenorgan slots, preserving rate and timing and destroying sender identity.

H.3 Learning

Agents are trained with recurrent multi-agent PPO [32, 33] under centralised training with decentralised execution: one shared actor (2-layer MLP encoder, LayerNorm, tanh, 128 units $\rightarrow$ GRU with 128 units [34])

and a centralised critic that observes all agents’ concatenated observations. The action space is hybrid—two Gaussians for thrust and turn, two Bernoullis for discharge and bite. Keeping the discharge decision explicitly Bernoulli matters for the analysis, because its probability is exactly the quantity we intervene on. GAE [35] with $\gamma = 0.99$, $\lambda = 0.95$; PPO clip 0.2; 4 epochs, 4 minibatches; Adam at 3×10^{-4} ; 512 parallel arenas; 64-step truncated BPTT; 18×10^6 agent-steps per run. Every condition is trained from scratch with independent seeds; *no seed selection of any kind is applied*, and all reported statistics pool over all seeds.

H.4 Measurement

Interventional causal influence. Jaques et al. [8] define the influence of agent j on agent i as the KL divergence between i ’s policy conditioned on j ’s actual action and its policy marginalised over j ’s alternatives, and estimate it with a learned model of other agents. Owing the simulator, we evaluate the counterfactual directly:

$$\text{CIE}_t^{j \rightarrow i} = D_{\text{KL}}[\pi_i(\cdot \mid h_t^i, o_t^i[\text{do}(e_j=1)]) \parallel \pi_i(\cdot \mid h_t^i, o_t^i[\text{do}(e_j=0)])], \quad (6)$$

where $o_t^i[\cdot]$ denotes re-rendering i ’s entire sensory input under the intervention with the recurrent state h_t^i and every other cause held on its factual path. The KL is analytic for the factored Gaussian $\times$ Bernoulli policy. The null is the same quantity computed between two re-renderings of the *same* intervention under independent sensory-noise draws: the divergence the policy shows when nothing about the message changed.

Positive signalling and positive listening. Positive signalling [7] is the mutual information between a window descriptor of the emitter’s pulse train (discharge count over 21 steps, quantised) and the emitter’s private state, normalised by the state entropy and computed *conditionally on* distance-to-nearest-neighbour decile. The stratification is essential: in a patchy world a fish near a neighbour is also near food, so an unconditioned mutual information can be manufactured entirely by geometry. Estimation is plug-in with Miller–Madow bias correction, with a permutation null. Positive listening is the L^1 divergence between the receiver’s factual policy and its policy under a time-shifted message, which preserves the marginal rate and destroys contingency—the null recommended by Lowe et al. [7] in preference to a silent channel, which is out of distribution.

Message content. What a receiver could recover is decoded *only* from receiver-observable quantities: the number of pulses it heard from that emitter in a trailing window and whether it heard anything. Ground-truth state is never given to the decoder. Targets are the number of live food items within 10 cm of the emitter, whether the emitter eats within the next 41 steps, emitter body size (dominance), emitter turn command (intended movement), and whether a predator is within 25 cm (danger). We report cross-validated R^2 or AUC minus the same quantity computed on time-shifted messages, so that any score attributable to marginal rate alone is subtracted off.

Sender Shaping Index. A signal, on Scott-Phillips’s definition, requires that production be shaped by its effect on receivers. We train matched populations in a *hearing* world and a *deaf* world (knollenorgans disabled throughout training) and compare their emission policies on a common, condition-neutral batch of observations:

$$\text{SSI} = \frac{\mathbb{E}_c D_{\text{JS}}[\pi_{\text{emit}}^{\mathcal{H}}(\cdot \mid c) \parallel \pi_{\text{emit}}^{\mathcal{D}}(\cdot \mid c)]}{\mathbb{E}_c D_{\text{JS}}[\pi_{\text{emit}}^{\text{seed } a} \parallel \pi_{\text{emit}}^{\text{seed } b}]} - 1. \quad (7)$$

The denominator is the within-world seed-to-seed divergence, so $\text{SSI} \leq 0$ means the emission policy differs no more between worlds than between seeds of the same world—i.e. nobody listening changes nothing, and the pulse is a *cue*. $\text{SSI} > 0$ means production is shaped by reception.

Payoff consequences and the honesty classification. Every channel condition is evaluated against the intact channel *paired within arena*: the two rollouts share the initial food layout and the sensory noise stream, so differencing removes essentially all layout variance. Confidence intervals are 2000-sample bootstraps over arenas. Comparing the intact channel with a rate-matched replay isolates the payoff value of the signal’s *contingency*. Following Maynard Smith and Harper [4] and Dawkins and Krebs [36], a condition is labelled

COMMUNICATION when contingency benefits both parties, MANIPULATION when it benefits the sender at the receiver’s expense, and EXPLOITED CUE when it benefits the receiver at the sender’s expense.

H.5 Validation against the reference implementation

Every stage of the sensory pipeline is checked against the original NumPy code in double precision on random scenes (`scripts/validate_physics.py`); results are in Table 2. Agreement is at floating-point round-off, and the knollenorgan and ampullary pathways agree bit-for-bit.

H.6 Deliberate divergences from the reference implementation

We record these explicitly because they bear on how far our results transfer.

1. **Wall-image amplitude.** Upstream’s configuration file sets a reflection coefficient of 0.95, but `ElectricScene..build.refl` never forwards it, so the code that produced the published results uses unattenuated images. We match the code, not the configuration file. Using 0.95 instead shifts the corollary-discharge baseline by $\sim 1.4 \times 10^{-6}$ relative and changes ampullary readings qualitatively, because those are a small difference of large numbers.
2. **Fixed arena and episode length.** Upstream samples arena dimensions per episode from $[40, 300]$ cm and does not report episode length. We fix a 60×60 cm arena and 512 steps so that conditions in a large factorial sweep are exactly comparable. Our conclusions are therefore about one arena scale.
3. **Collision resolution** is synchronous (all fish propose, then all resolve) rather than in agent order.
4. **Network size.** 128-unit GRU rather than 512, and 18×10^6 agent-steps per run rather than 5×10^6 environment steps, chosen to fit 354 independent training runs into the compute budget.
5. **Sensory noise** is present (multiplicative, fractions given above), following the upstream code; the upstream paper text describes a noiseless model.
6. **No seed selection.** Upstream selects a single seed (S1) by an ethological criterion and reports analyses on it. We pool over all seeds.